

# CHANGES: Findora (after the poster session)

---

Feedback came from **6 peer reviewers**. Below we mention every piece of feedback, what we changed in response and why — and, where we choose **not** to change something, the reasoning. The post-change **evaluation** is at the end.

---

## 1. What we changed

"Expand the data: add product details, descriptions and reviews."  
(reviewer 2; also our own stated limitation)

**\*\*Changed:** \*\* Every phone now has a grounded **description** and a clearly-labeled **AI review summary**, generated *only* from its real specs + rating (`enrich_smartphones.py` → `datasets/smartphone_enrichment.json`, joined into the catalog at load time, similar to image cache). They appear in a product's **Details** view and feed the semantic "vibe" search. Coverage is **100% of the 500 phones**, and the summary is explicitly marked "*AI summary (generated from specs & rating)*" so it is never mistaken for a real user review.

*Why generated rather than scraped real reviews:* we made an attempt to get the reviews for the matching products but could not find the reviews for all the products in one dataset.

"Show one more real-world example of the full recommendation process." (reviewer 7)

**Changed:** Added in the updated poster.

"Address the local-model limitation and the agent's failure handling more concretely." (reviewer 6)

**Changed:** Bad tool arguments are validated and **never crash** (tools never raise); a request needing more than the 4-round budget is forced to a final grounded answer; local function-calling is unreliable on small models, which is *why* the system is cloud-primary.

"Be honest about / improve loading times." (reviewer 4)

**Changed:** We tried through prompt caching but still depends on 2 API calls.

"Surface the less-intuitive commands (e.g. you can tell it to clear all filters)." (reviewer 4)

**Changed:** Added a "**Handy commands**" section to the in-app " *How it works*" panel listing the non-obvious natural-language commands — "*clear all filters*", "*ignore that*" (*undo*), "*cheaper ones*", "*compare those*", "*why this one?*", "*actually, show me headphones*" — so users discover them.

"The '700 True' tooltip was confusing." (reviewer 5)

**Resolved:** We could not reproduce it in the current build: the post-session UI rework relabeled every spec/score value with a unit or **Yes/No** (verified across the spec table, the score breakdown, and the compare

view), so no raw value renders without a label. It appears the rework already eliminated it.

## 2. Considered but not changed (with reasoning)

- **Cross-shop price comparison:** Out of scope: the reviewer agreed. Findora is a single-catalog recommender; live multi-retailer pricing needs commercial price-feed integrations well beyond this project.
  - **Expanding the catalog to ~1,000 phones:** We deliberately **enriched the existing 500 in place** rather than swap the dataset: it keeps our validated evaluation baseline comparable, keeps the clean USD prices the ranking depends on, and the enrichment (not the row count) was the actual gap.
  - **Lower loading time via response streaming:** Acknowledged, not implemented. The agent is already faster on average and we added the honest latency analysis; streaming the final reply (lower *perceived* latency) is the clear next step but was deprioritised in favour of the data work.
-